

My Details

Name : Rohan Ingle

PRN : 22070126047

Batch : AIML A2 (2022-2026)

GitHub Link : <https://github.com/Rohan-ingle/Natural-Language-Processing>

Importing Required Libraries

In [2]:

```
import warnings

import re
import pandas as pd
import cudf
import cupy as cp

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import classification_report, confusion_matrix

from scipy.sparse import csr_matrix
from cuml.svm import SVC as cuSVC
from cuml.ensemble import RandomForestClassifier as cuRF

import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer

import seaborn as sns
import matplotlib.pyplot as plt

warnings.filterwarnings("ignore")
```

In [2]:

```
nltk.download('all')
```

```
[nltk_data] Downloading collection 'all'
[nltk_data] |
[nltk_data] | Downloading package abc to
[nltk_data] |   /home/speedindeed1/nltk_data...
[nltk_data] | Package abc is already up-to-date!
[nltk_data] | Downloading package alpino to
[nltk_data] |   /home/speedindeed1/nltk_data...
[nltk_data] | Package alpino is already up-to-date!
[nltk_data] | Downloading package averaged_perceptron_tagger to
[nltk_data] |   /home/speedindeed1/nltk_data...
[nltk_data] | Package averaged_perceptron_tagger is already up-
[nltk_data] | to-date!
[nltk_data] | Downloading package averaged_perceptron_tagger_eng to
[nltk_data] |   /home/speedindeed1/nltk_data...
[nltk_data] | Package averaged_perceptron_tagger_eng is already
[nltk_data] | up-to-date!
[nltk_data] | Downloading package averaged_perceptron_tagger_ru to
[nltk_data] |   /home/speedindeed1/nltk_data...
[nltk_data] | Package averaged_perceptron_tagger_ru is already
[nltk_data] | up-to-date!
[nltk_data] | Downloading package averaged_perceptron_tagger_rus to
[nltk_data] |   /home/speedindeed1/nltk_data...
[nltk_data] | Package averaged_perceptron_tagger_rus is already
[nltk_data] | up-to-date!
[nltk_data] | Downloading package basque_grammars to
[nltk_data] |   /home/speedindeed1/nltk_data...
[nltk_data] | Package basque_grammars is already up-to-date!
[nltk_data] | Downloading package bcp47 to
[nltk_data] |   /home/speedindeed1/nltk_data...
[nltk_data] | Package bcp47 is already up-to-date!
[nltk_data] | Downloading package biocreative_ppi to
[nltk_data] |   /home/speedindeed1/nltk_data...
[nltk_data] | Package biocreative_ppi is already up-to-date!
```

```
[nltk_data] | Downloading package blip_wsj_no_aux to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package blip_wsj_no_aux is already up-to-date!
[nltk_data] | Downloading package book_grammars to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package book_grammars is already up-to-date!
[nltk_data] | Downloading package brown to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package brown is already up-to-date!
[nltk_data] | Downloading package brown_tei to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package brown_tei is already up-to-date!
[nltk_data] | Downloading package cess_cat to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package cess_cat is already up-to-date!
[nltk_data] | Downloading package cess_esp to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package cess_esp is already up-to-date!
[nltk_data] | Downloading package chat80 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package chat80 is already up-to-date!
[nltk_data] | Downloading package city_database to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package city_database is already up-to-date!
[nltk_data] | Downloading package cmudict to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package cmudict is already up-to-date!
[nltk_data] | Downloading package comparative_sentences to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package comparative_sentences is already up-to-
[nltk_data] | date!
[nltk_data] | Downloading package comtrans to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package comtrans is already up-to-date!
[nltk_data] | Downloading package conll2000 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package conll2000 is already up-to-date!
[nltk_data] | Downloading package conll2002 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package conll2002 is already up-to-date!
[nltk_data] | Downloading package conll2007 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package conll2007 is already up-to-date!
[nltk_data] | Downloading package crubadan to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package crubadan is already up-to-date!
[nltk_data] | Downloading package dependency_treebank to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package dependency_treebank is already up-to-date!
[nltk_data] | Downloading package dolch to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package dolch is already up-to-date!
[nltk_data] | Downloading package europarl_raw to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package europarl_raw is already up-to-date!
[nltk_data] | Downloading package extended_omw to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package extended_omw is already up-to-date!
[nltk_data] | Downloading package floresta to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package floresta is already up-to-date!
[nltk_data] | Downloading package framenet_v15 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package framenet_v15 is already up-to-date!
[nltk_data] | Downloading package framenet_v17 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package framenet_v17 is already up-to-date!
[nltk_data] | Downloading package gazetteers to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package gazetteers is already up-to-date!
[nltk_data] | Downloading package genesis to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package genesis is already up-to-date!
[nltk_data] | Downloading package gutenberg to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package gutenberg is already up-to-date!
[nltk_data] | Downloading package ier to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package ier is already up-to-date!
[nltk_data] | Downloading package inaugural to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package inaugural is already up-to-date!
[nltk_data] | Downloading package indian to
```

```
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package indian is already up-to-date!
[nltk_data] | Downloading package jeita to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package jeita is already up-to-date!
[nltk_data] | Downloading package kimmo to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package kimmo is already up-to-date!
[nltk_data] | Downloading package knbc to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package knbc is already up-to-date!
[nltk_data] | Downloading package large_grammars to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package large_grammars is already up-to-date!
[nltk_data] | Downloading package lin_thesaurus to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package lin_thesaurus is already up-to-date!
[nltk_data] | Downloading package mac_morpho to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package mac_morpho is already up-to-date!
[nltk_data] | Downloading package machado to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package machado is already up-to-date!
[nltk_data] | Downloading package masc_tagged to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package masc_tagged is already up-to-date!
[nltk_data] | Downloading package maxent_ne_chunker to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package maxent_ne_chunker is already up-to-date!
[nltk_data] | Downloading package maxent_ne_chunker_tab to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package maxent_ne_chunker_tab is already up-to-
[nltk_data] | date!
[nltk_data] | Downloading package maxent_treebank_pos_tagger to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package maxent_treebank_pos_tagger is already up-
[nltk_data] | to-date!
[nltk_data] | Downloading package maxent_treebank_pos_tagger_tab to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package maxent_treebank_pos_tagger_tab is already
[nltk_data] | up-to-date!
[nltk_data] | Downloading package mooses_sample to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package mooses_sample is already up-to-date!
[nltk_data] | Downloading package movie_reviews to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package movie_reviews is already up-to-date!
[nltk_data] | Downloading package mte_teip5 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package mte_teip5 is already up-to-date!
[nltk_data] | Downloading package mwa_ppdb to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package mwa_ppdb is already up-to-date!
[nltk_data] | Downloading package names to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package names is already up-to-date!
[nltk_data] | Downloading package nombank.1.0 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package nombank.1.0 is already up-to-date!
[nltk_data] | Downloading package nonbreaking_prefixes to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package nonbreaking_prefixes is already up-to-date!
[nltk_data] | Downloading package nps_chat to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package nps_chat is already up-to-date!
[nltk_data] | Downloading package omw to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package omw is already up-to-date!
[nltk_data] | Downloading package omw-1.4 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package omw-1.4 is already up-to-date!
[nltk_data] | Downloading package opinion_lexicon to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package opinion_lexicon is already up-to-date!
[nltk_data] | Downloading package panlex_swadesh to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package panlex_swadesh is already up-to-date!
[nltk_data] | Downloading package paradigms to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package paradigms is already up-to-date!
[nltk_data] | Downloading package pe08 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package pe08 is already up-to-date!
```

```
[nltk_data] | Downloading package perluniprops to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package perluniprops is already up-to-date!
[nltk_data] | Downloading package pil to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package pil is already up-to-date!
[nltk_data] | Downloading package pl196x to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package pl196x is already up-to-date!
[nltk_data] | Downloading package porter_test to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package porter_test is already up-to-date!
[nltk_data] | Downloading package ppattach to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package ppattach is already up-to-date!
[nltk_data] | Downloading package problem_reports to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package problem_reports is already up-to-date!
[nltk_data] | Downloading package product_reviews_1 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package product_reviews_1 is already up-to-date!
[nltk_data] | Downloading package product_reviews_2 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package product_reviews_2 is already up-to-date!
[nltk_data] | Downloading package propbank to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package propbank is already up-to-date!
[nltk_data] | Downloading package pros_cons to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package pros_cons is already up-to-date!
[nltk_data] | Downloading package ptb to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package ptb is already up-to-date!
[nltk_data] | Downloading package punkt to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package punkt is already up-to-date!
[nltk_data] | Downloading package punkt_tab to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package punkt_tab is already up-to-date!
[nltk_data] | Downloading package qc to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package qc is already up-to-date!
[nltk_data] | Downloading package reuters to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package reuters is already up-to-date!
[nltk_data] | Downloading package rslp to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package rslp is already up-to-date!
[nltk_data] | Downloading package rte to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package rte is already up-to-date!
[nltk_data] | Downloading package sample_grammars to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package sample_grammars is already up-to-date!
[nltk_data] | Downloading package semcor to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package semcor is already up-to-date!
[nltk_data] | Downloading package senseval to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package senseval is already up-to-date!
[nltk_data] | Downloading package sentence_polarity to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package sentence_polarity is already up-to-date!
[nltk_data] | Downloading package sentiwordnet to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package sentiwordnet is already up-to-date!
[nltk_data] | Downloading package shakespeare to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package shakespeare is already up-to-date!
[nltk_data] | Downloading package sinica_treebank to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package sinica_treebank is already up-to-date!
[nltk_data] | Downloading package smultron to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package smultron is already up-to-date!
[nltk_data] | Downloading package snowball_data to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package snowball_data is already up-to-date!
[nltk_data] | Downloading package spanish_grammars to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package spanish_grammars is already up-to-date!
[nltk_data] | Downloading package state_union to
[nltk_data] | /home/speedindeed1/nltk_data...
```

```
[nltk_data] | Package state_union is already up-to-date!
[nltk_data] | Downloading package stopwords to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package stopwords is already up-to-date!
[nltk_data] | Downloading package subjectivity to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package subjectivity is already up-to-date!
[nltk_data] | Downloading package swadesh to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package swadesh is already up-to-date!
[nltk_data] | Downloading package switchboard to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package switchboard is already up-to-date!
[nltk_data] | Downloading package tagsets to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package tagsets is already up-to-date!
[nltk_data] | Downloading package tagsets_json to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package tagsets_json is already up-to-date!
[nltk_data] | Downloading package timit to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package timit is already up-to-date!
[nltk_data] | Downloading package toolbox to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package toolbox is already up-to-date!
[nltk_data] | Downloading package treebank to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package treebank is already up-to-date!
[nltk_data] | Downloading package twitter_samples to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package twitter_samples is already up-to-date!
[nltk_data] | Downloading package udhr to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package udhr is already up-to-date!
[nltk_data] | Downloading package udhr2 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package udhr2 is already up-to-date!
[nltk_data] | Downloading package unicode_samples to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package unicode_samples is already up-to-date!
[nltk_data] | Downloading package universal_tagset to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package universal_tagset is already up-to-date!
[nltk_data] | Downloading package universal_treebanks_v20 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package universal_treebanks_v20 is already up-to-
[nltk_data] | date!
[nltk_data] | Downloading package vader_lexicon to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package vader_lexicon is already up-to-date!
[nltk_data] | Downloading package verbnet to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package verbnet is already up-to-date!
[nltk_data] | Downloading package verbnet3 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package verbnet3 is already up-to-date!
[nltk_data] | Downloading package webtext to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package webtext is already up-to-date!
[nltk_data] | Downloading package wmt15_eval to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package wmt15_eval is already up-to-date!
[nltk_data] | Downloading package word2vec_sample to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package word2vec_sample is already up-to-date!
[nltk_data] | Downloading package wordnet to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package wordnet is already up-to-date!
[nltk_data] | Downloading package wordnet2021 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package wordnet2021 is already up-to-date!
[nltk_data] | Downloading package wordnet2022 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package wordnet2022 is already up-to-date!
[nltk_data] | Downloading package wordnet31 to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package wordnet31 is already up-to-date!
[nltk_data] | Downloading package wordnet_ic to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package wordnet_ic is already up-to-date!
[nltk_data] | Downloading package words to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package words is already up-to-date!
```

```
[nltk_data] | Downloading package ycoe to
[nltk_data] | /home/speedindeed1/nltk_data...
[nltk_data] | Package ycoe is already up-to-date!
[nltk_data] |
[nltk_data] Done downloading collection all
```

```
Out[2]: True
```

Loading Datasets into a DataFrame

```
In [3]: train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')
```

Basic Preprocessing

- Lets select only necessary features

```
In [4]: train["Genre"].unique()
```

```
Out[4]: array(['Rock', 'Metal', 'Pop', 'Indie', 'R&B', 'Folk', 'Electronic',
              'Jazz', 'Hip-Hop', 'Country'], dtype=object)
```

```
In [5]: test["Genre"].unique()
```

```
Out[5]: array(['Hip-Hop', 'Indie', 'Metal', 'Pop', 'Country', 'Jazz', 'Rock',
              'R&B', 'Electronic', 'Folk'], dtype=object)
```

Lets rename **Hip-Hop** to **HH** for simplicity

```
In [4]: train['Genre'] = train['Genre'].str.replace('Hip-Hop', 'HH')
test['Genre'] = test['Genre'].str.replace('Hip-Hop', 'HH')
```

```
In [7]: train["Genre"].unique()
```

```
Out[7]: array(['Rock', 'Metal', 'Pop', 'Indie', 'R&B', 'Folk', 'Electronic',
              'Jazz', 'HH', 'Country'], dtype=object)
```

```
In [8]: test["Genre"].unique()
```

```
Out[8]: array(['HH', 'Indie', 'Metal', 'Pop', 'Country', 'Jazz', 'Rock', 'R&B',
              'Electronic', 'Folk'], dtype=object)
```

Here we will only select **English** i.e. **'en'** songs only

```
In [5]: genre = ['Rock', 'Jazz', 'HH', 'Metal', 'Country']
train = train[(train['Genre'].isin(genre)) & (train['Language'] == 'en')]
```

```
In [10]: train.describe()
```

```
Out[10]:
```

	Artist	Song	Genre	Language	Lyrics
count	143720	143719	143720	143720	143720
unique	6193	86662	5	1	127917
top	elvis presley	home	Rock	en	My heart is sad and lonely\nFor you I sigh, fo...
freq	1598	87	107145	143720	83

Get insights on duplicated values

```
In [6]: duplicates = train[train['Lyrics'].duplicated(keep=False)].sort_values(by='Lyrics')
```

```
In [12]: duplicates
```

Out [12]:

	Artist	Song	Genre	Language	Lyrics
12587	cowboy junkies	take me	Rock	en	\n\nwhere the thunder meets the light\nTake me...
12469	cowboy junkies	take me	Rock	en	\n\nwhere the thunder meets the light\nTake me...
30	12 stones	back up	Rock	en	\nI hear the words you say to me\nI see the wa...
9	12 stones	back up	Rock	en	\nI hear the words you say to me\nI see the wa...
245656	george strait	lovesick blues	Rock	en	\nI'm in love,\nI'm in love,\nI'm in love with...
...
22528	george michael	you've changed	Rock	en	you've changed\nthat sparkle in your eyes has ...
9402	bush	dead meat	Rock	en	your dead meat\nyour dead meat\nyour dead meat...
9431	bush	dead meat	Rock	en	your dead meat\nyour dead meat\nyour dead meat...
11969	coheed and cambria	a favor house atlantic	Metal	en	your eyes tell the stories of a day you wish y...
11963	coheed and cambria	a favor house atlantic	Metal	en	your eyes tell the stories of a day you wish y...

26207 rows × 5 columns

In [13]:

```
duplicates_by_genre = train[train.duplicated(subset=['Lyrics', 'Genre'], keep=False)]

duplicate_counts_by_genre = duplicates_by_genre.groupby('Genre').size()

print(duplicate_counts_by_genre)
```

Genre
Jazz 6020
Metal 2495
Rock 16985
dtype: int64

In [14]:

```
train["Genre"].unique()
```

Out[14]: array(['Rock', 'Metal', 'Jazz', 'HH', 'Country'], dtype=object)

In [15]:

```
test["Genre"].unique()
```

Out[15]: array(['HH', 'Indie', 'Metal', 'Pop', 'Country', 'Jazz', 'Rock', 'R&B',
'Electronic', 'Folk'], dtype=object)

In [16]:

```
test.describe()
```

Out[16]:

	Song year	Track_id
count	7935.000000	7935.000000
mean	2007.943793	13262.260744
std	5.201110	7612.173049
min	1970.000000	3.000000
25%	2006.000000	6708.500000
50%	2007.000000	13169.000000
75%	2011.000000	19840.000000
max	2016.000000	26446.000000

lets define our `stop_words` and `lemmatizer`

In [7]:

```
stop_words = set(stopwords.words('english'))
lemmatizer = WordNetLemmatizer()
```

Here we will make a preprocessing function to clean the lyrics string and also tokenize it

In [8]:

```
def text_cleaner(text):
    text = re.sub(r'\@w+|\#','', text)
    text = re.sub(r'^A-Za-z\s','', text)
    text = text.lower()
    tokens = word_tokenize(text)
    tokens = [lemmatizer.lemmatize(word) for word in tokens if word not in stop_words]
    return ' '.join(tokens)
```

In [9]:

```
train['Lyrics'] = train['Lyrics'].apply(text_cleaner)
```

In [10]:

```
test = test[(test['Genre'].isin(genre))]
test['Lyrics'] = test['Lyrics'].apply(text_cleaner)
```

Lets check how it affected **Lyrics** column

```
In [21]: train.head()
```

Out[21]:	Artist	Song	Genre	Language	Lyrics
0	12 stones	world so cold	Rock	en	start pain followed hate fueled endless questi...
1	12 stones	broken	Rock	en	freedom alone alone patiently waiting phone ho...
2	12 stones	3 leaf loser	Rock	en	biting hand feed lying voice inside reach beg ...
3	12 stones	anthem for the underdog	Rock	en	say know cant imagine wait across line thought...
4	12 stones	adrenaline	Rock	en	heart beating faster cant control feeling anym...

```
In [22]: test.head()
```

Out[22]:		Song	Song year	Artist	Genre	Lyrics	Track_id
0		craftsmanship	2005	buck-65	HH	folk spend day daydreaming finding clue whole ...	8294
2		riot	2013	bullet-for-my-valentine	Metal	ready time war well break fucking door smash w...	3301
4		believe-in-a-dollar	2012	cassidy	HH	believe magic young girl heart music free when...	16797
5		mama-bake-a-pie-daddy-kill-a-chicken	2007	bobby-bare	Country	people starin wheel ramp toward plane war ive ...	1251
6		thinking-about-you	2007	bill-monroe	Country	sigh hour away think love yesterday know shes ...	25217

Now we are going to encode our dataset, define vectorizers and models and transform data by vectorizing it

- We will use `fit_transform` on training dataset and `transform` on testing

```
In [11]: encode_label = LabelEncoder()  
train['Genre'] = encode_label.fit_transform(train['Genre'])  
test['Genre'] = encode_label.transform(test['Genre'])
```

```
In [12]: countVec = CountVectorizer(max_features=50000)  
tfidfVec = TfidfVectorizer(max_features=50000)
```

```
In [13]: X_train_countVec = countVec.fit_transform(train['Lyrics'])  
X_test_countVec = countVec.transform(test['Lyrics'])  
  
X_train_tfidfVec = tfidfVec.fit_transform(train['Lyrics'])  
X_test_tfidfVec = tfidfVec.transform(test['Lyrics'])
```

```
In [14]: y_train = train['Genre']  
y_test = test['Genre']
```

Training


We will make a function to train models

```
In [15]: def train(model, X_train, X_test, y_train, y_test, model_name, vectorizer):  
    model.fit(X_train, y_train)  
    y_pred = model.predict(X_test)  
  
    print(f"\nClassification Report for {model_name} with {vectorizer}:")  
    print(classification_report(y_test, y_pred, target_names=encode_label.classes_))  
  
    cm = confusion_matrix(y_test, y_pred)  
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=encode_label.classes_, yticklabels=encode_la  
    plt.title(f'Confusion Matrix for {model_name} with {vectorizer}')  
    plt.xlabel('Predicted')  
    plt.ylabel('True')  
    plt.show()
```

```
In [28]: lr = LogisticRegression(class_weight='balanced', max_iter=10000)
```


```
In [29]: train(lr, X_train_tfidfVec, X_test_tfidfVec, y_train, y_test, 'logisticRegression', 'Tfidf Vectorizer')
```


Classification Report for logisticRegression with Tfidf Vectorizer:				
	precision	recall	f1-score	support
Country	0.64	0.47	0.54	810
HH	0.92	0.78	0.84	960
Jazz	0.58	0.52	0.55	660
Metal	0.59	0.83	0.69	810
Rock	0.52	0.55	0.54	1410
accuracy			0.63	4650
macro avg	0.65	0.63	0.63	4650
weighted avg	0.64	0.63	0.63	4650



```
In [30]: train(lr, X_train_countVec, X_test_countVec, y_train, y_test,"Logistic Regression", 'Count Vectorizer')
```

Classification Report for Logistic Regression with Count Vectorizer:				
	precision	recall	f1-score	support
Country	0.66	0.28	0.39	810
HH	0.93	0.59	0.72	960
Jazz	0.56	0.54	0.55	660
Metal	0.60	0.74	0.66	810
Rock	0.46	0.67	0.54	1410
accuracy			0.58	4650
macro avg	0.64	0.56	0.57	4650
weighted avg	0.63	0.58	0.57	4650



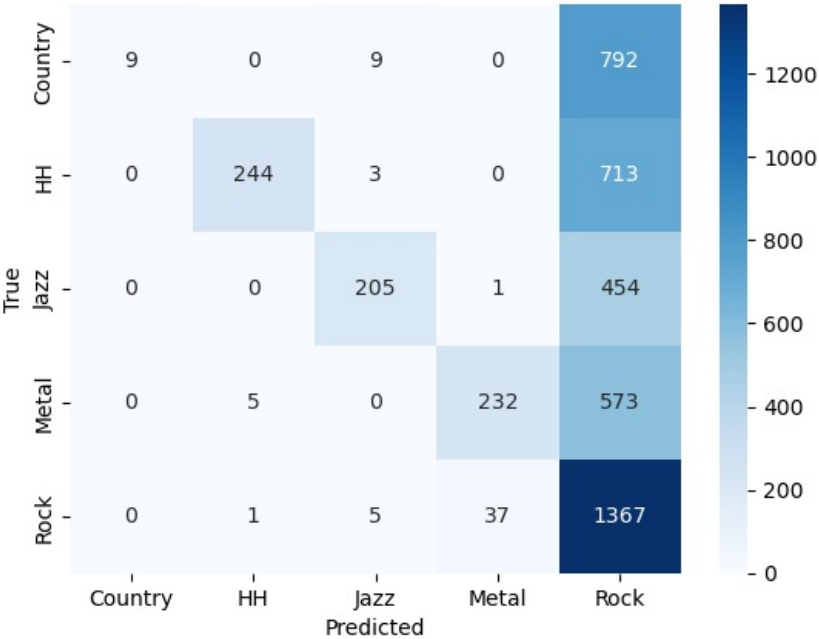
```
In [31]: rf = RandomForestClassifier(class_weight='balanced')
```

```
In [32]: train(rf, X_train_tfidfVec, X_test_tfidfVec, y_train, y_test, 'Random Forest', 'Tfidf Vectorizer')
```

Classification Report for Random Forest with Tfidf Vectorizer:

	precision	recall	f1-score	support
Country	1.00	0.01	0.02	810
HH	0.98	0.25	0.40	960
Jazz	0.92	0.31	0.46	660
Metal	0.86	0.29	0.43	810
Rock	0.35	0.97	0.51	1410
accuracy			0.44	4650
macro avg	0.82	0.37	0.37	4650
weighted avg	0.76	0.44	0.38	4650

Confusion Matrix for Random Forest with Tfidf Vectorizer




```
In [33]: train(rf, X_train_countVec, X_test_countVec, y_train, y_test, 'Random Forest', 'Count Vectorizer')
```

Classification Report for Random Forest with Count Vectorizer:

	precision	recall	f1-score	support
Country	0.88	0.02	0.03	810
HH	0.97	0.25	0.39	960
Jazz	0.92	0.32	0.47	660
Metal	0.85	0.28	0.42	810
Rock	0.35	0.97	0.52	1410
accuracy			0.44	4650
macro avg	0.79	0.37	0.37	4650
weighted avg	0.74	0.44	0.38	4650

Confusion Matrix for Random Forest with Count Vectorizer



Using cuML to train SVC on GPU to improve training time

But first we need to convert data to support cuML library and models

- Converting to `csr_matrix` ensures that sparse data, like word vectors from `CountVectorizer` or `TF-IDF`, is stored efficiently by keeping only non-zero values. This reduces memory usage and speeds up operations on large datasets.
- Switching to `float32` cuts memory usage in half compared to `float64`. This format is more compatible with GPU-accelerated libraries like `CuPy`, offering a good balance between precision and performance.
- `CuPy` allows data to be processed on NVIDIA GPUs, speeding up computation for large datasets and complex operations. By converting to `CuPy`'s sparse format, you take full advantage of parallel GPU processing for faster model training.

```
In [16]: X_train_count_float32_csrMatrix = csr_matrix(X_train_countVec, dtype=cp.float32)
X_test_count_float32_csrMatrix = csr_matrix(X_test_countVec, dtype=cp.float32)
```

```
In [ ]: X_train_tfidf_float32_csrMatrix = csr_matrix(X_train_tfidfVec, dtype=cp.float32)
X_test_tfidf_float32_csrMatrix = csr_matrix(X_test_tfidfVec, dtype=cp.float32)
```

```
In [ ]: X_train_count_cupy = cp.sparse.csr_matrix(X_train_count_float32_csrMatrix)
X_test_count_cupy = cp.sparse.csr_matrix(X_test_count_float32_csrMatrix)
```

```
In [ ]: X_train_tfidf_cupy = cp.sparse.csr_matrix(X_train_tfidf_float32_csrMatrix)
X_test_tfidf_cupy = cp.sparse.csr_matrix(X_test_tfidf_float32_csrMatrix)
```

```
In [ ]: y_train_cupy = cp.array(y_train.to_numpy(), dtype=cp.int32)
y_test_cupy = cp.array(y_test.to_numpy(), dtype=cp.int32)
```

Here we will make a function to train SVC models using cuML

```
In [17]: def cuML_train(model, X_train, X_test, y_train, y_test, model_name, vectorizer_name):
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)

        print(f"\nClassification Report for {model_name} with {vectorizer_name}:")
        print(classification_report(cp.asnumpy(y_test), cp.asnumpy(y_pred), target_names=encode_label.classes_))

        confusionMatrix = confusion_matrix(cp.asnumpy(y_test), cp.asnumpy(y_pred))
        sns.heatmap(confusionMatrix, annot=True, fmt='d', cmap='Blues', xticklabels=encode_label.classes_, yticklabels=encode_label.classes_)
        plt.title(f'Confusion Matrix for {model_name} with {vectorizer_name}')
        plt.xlabel('Predicted')
        plt.ylabel('True')
        plt.show()
```

Defining model

```
In [18]: SVC_cuML = cuSVC(kernel = "rbf", probability=True, class_weight='balanced')
```


```
In [37]: cuML_train(SVC_cuML, X_train_count_cupy, X_test_count_cupy, y_train_cupy, y_test_cupy, 'cuML SVC', 'Count Vectorizer')
```

[W] [05:00:58.749077] Sample weights are currently ignored for multi class classification
[W] [05:02:38.408955] Sample weights are currently ignored for multi class classification
[W] [05:04:18.159253] Sample weights are currently ignored for multi class classification
[W] [05:05:58.520177] Sample weights are currently ignored for multi class classification
[W] [05:07:37.618347] Sample weights are currently ignored for multi class classification

Classification Report for cuML SVC with Count Vectorizer:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Country	0.56	0.56	0.56	810
HH	0.82	0.86	0.84	960
Jazz	0.58	0.56	0.57	660
Metal	0.63	0.82	0.71	810
Rock	0.57	0.46	0.51	1410
accuracy			0.64	4650
macro avg	0.63	0.65	0.64	4650
weighted avg	0.63	0.64	0.63	4650




```
In [19]: cuML_train(SVC_cuML, X_train_tfidf_cupy, X_test_tfidf_cupy, y_train_cupy, y_test_cupy, 'cuML SVC', 'Tfidf Vectorizer')
```

[W] [14:12:26.733111] Sample weights are currently ignored for multi class classification
[W] [14:14:20.595877] Sample weights are currently ignored for multi class classification
[W] [14:16:15.122395] Sample weights are currently ignored for multi class classification
[W] [14:18:06.489816] Sample weights are currently ignored for multi class classification
[W] [15:32:02.664048] Sample weights are currently ignored for multi class classification

Classification Report for cuML SVC with Tfidf Vectorizer:

	precision	recall	f1-score	support
Country	0.62	0.57	0.60	810
HH	0.89	0.80	0.85	960
Jazz	0.77	0.52	0.62	660
Metal	0.70	0.78	0.74	810
Rock	0.55	0.66	0.60	1410
accuracy			0.67	4650
macro avg	0.71	0.67	0.68	4650
weighted avg	0.69	0.67	0.68	4650



Insights

1. Logistic Regression with Count Vectorizer and Tf-idf Vectorizer:

Count Vectorizer:

- **Precision, recall, and F1-scores** are low for several classes, especially for "Country" and "Rock" genres, indicating difficulty in predicting these classes.
- **Confusion Matrix** shows a significant amount of misclassification for these classes. For example, many "Country" predictions are classified as "Rock."

Tf-idf Vectorizer:

- **Slight improvement in performance** over Count Vectorizer, especially for "Rock" and "Jazz."
- **Confusion Matrix** shows fewer misclassifications for some classes (e.g., "Rock"), but the model still struggles with distinguishing "Country" and "Jazz."

2. Random Forest with Count Vectorizer and Tf-idf Vectorizer:

Count Vectorizer:

- **Overall performance is poor**, particularly for "Country" and "Jazz."
- **Confusion Matrix** indicates the model is skewed towards predicting "Rock," with many other classes misclassified as "Rock."

Tf-idf Vectorizer:

- **Similar issues as Count Vectorizer**, though there is a slight improvement in performance metrics.
- **Confusion Matrix** shows improvement in predicting "HH" and "Jazz," but "Country" is still largely misclassified as "Rock."

3. cuML SVC with Count Vectorizer and Tf-idf Vectorizer (GPU-accelerated):

Count Vectorizer:

- **Significant improvement in performance**, especially in recall and F1-scores for most classes compared to previous models.
- **Confusion Matrix** shows better prediction of "HH" and "Jazz" with fewer misclassifications.

Tf-idf Vectorizer:

- **Best performance among all models**, with higher precision and recall scores across all classes.
- **Confusion Matrix** reflects fewer misclassifications, particularly for "Rock" and "Jazz," making this model the most accurate.

Key Insights:

- **Tf-idf Vectorizer** generally performs better than **Count Vectorizer** across all models.
- **cuML SVC** offers the best performance due to GPU acceleration, providing faster and more accurate predictions compared to Logistic Regression and Random Forest.
- There is still room for improvement in distinguishing between certain genres like "Country" and "Rock."

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js